

Vector Space Model

Introduction

COMP3009J: Information Retrieval

Dr. David Lillis (david.lillis@ucd.ie)

UCD School of Computer Science
Beijing Dublin International College

Introduction

- So far, we have seen Information Retrieval (IR) as a process where a user submits a **query** consisting of some **terms** connected with **boolean operators**.
- This is known as the “Boolean Model” of IR, and has some serious disadvantages:
 - The model predicts that each document is either **relevant** or **non-relevant** to the **query**.
 - There is no notion of a “partial” match to the query conditions and so this can lead to too few documents being retrieved.
 - Every document that matches the query is treated the same, so **no ranking** occurs.

Introduction

- One of the most common IR models in use is the **Vector Space Model**.
- This models both **documents** *and* **queries** in an **N-dimensional space**.
- Basic vector algebra is used to calculate the **most similar** documents to a query.
- The basic vector space model has been adapted for many other purposes.
- The most common are:
 - Machine Learning (e.g. K-NN algorithm)
 - Clustering
- See page 77 of Modern Information Retrieval (2nd Ed.).

Basis

- A collection contains a set of documents each of which is composed of a “bag of words” (also known as terms).
 - Each unique **term in the collection** is represented by a dimension in the vector space.
 - A document is then represented by giving a non-zero value to the dimension representing a term it contains.
 - This number is known as a **term weight**.
 - Other dimensions get a value of zero.
 - i.e. the weight of a term is 0 in documents that do not contain it.
 - Note the order of the terms within the document is not retained.
- A **similarity score** is calculated based on the proximity of two vectors in the N-dimensional space.

Partial Matching

- Unlike the Boolean Model, the Vector Space Model facilitates **partial matching** by using **non-binary** term weights.
- In models like this, each document must have a **similarity score** calculated for it, which measures how similar it is to the given query.
- This is usually shown as $sim(q, d)$ (i.e. the similarity between a query q and a document d).
- These models return a **ranked list** of documents, where the documents with the highest similarity scores are at the top of the list.
- In this way, it is hoped that the most relevant documents are at the top of the result set, so that the user can find relevant information easier.

Vector Space Model

Document and Query Representation

Document Representation

- Consider the following documents:
 1. Information Retrieval is an exciting subject
 2. Mathematics is important in Information Retrieval
- Ignoring stopwords the collection contains **6 distinct terms**: information; retrieval; exciting; subject; mathematics; important.
- The vector space will have **six dimensions** (one for each term in the document collection).
- This means that **EVERY DOCUMENT** is mathematically represented by six dimensions, regardless of how many terms it contains.

Document Representation

- For the introduction to this topic, we first assume a simple weighting scheme (we will replace this later):
 - $w_{i,j}$ refers to the weight of term i in document j .
 - if term i is contained in document j , the weight is 1 (i.e. $w_{i,j} = 1$).
 - if term i is not contained in document j , $w_{i,j} = 0$.
- In practice, the weight will be a real number where:
 - $0 \leq w_{i,j} \leq 1$
- We will see how this is calculated later.

Term	Doc 1	Doc 2
information	1	1
retrieval	1	1
exciting	1	0
subject	1	0
mathematics	0	1
important	0	1

Query Representation

- The query is represented in an identical way.
- Consider the query: *important information*.
- This is represented in the vector model as:
 - $(1, 0, 0, 0, 0, 1)$
- It is important that the **order of the terms** in the **vector** is the same as for the corpus (i.e. the first dimension relates to the term “information”, the second relates to “retrieval”, etc.).
 - If the order changes, then we are no longer comparing like with like.

Vector Space Model

- The rest of this topic is organised as follows:
 - **Cosine Similarity:** how to compare queries with documents to measure how similar they are.
 - **Term Weights:** Some words are more important in describing a document than others.
 - **Conclusions and Summary**